

# Carlos Rivera

[email] carlos@example.com    [phone] +1 780-555-0144    [li] linkedin.com/in/carlosrivera    [gh] github.com/carlosrivera

## Summary

Data engineer with 6 years of experience designing ETL pipelines, data warehouses, and ML feature stores. Strong background in Python, Spark, AWS, and PostgreSQL.

## Technical Skills

|            |              |              |
|------------|--------------|--------------|
| Python     | Apache Spark | AWS Redshift |
| PostgreSQL | Docker       | Kubernetes   |
| LLMs       | RAG          | LangChain    |
| FastAPI    | CI/CD        | Airflow      |

## Professional Experience

**Senior Data Engineer**, DataCo Apr 2019 – Present  
Edmonton, AB

- Designed Spark-based ETL pipelines processing 5TB/day into AWS Redshift.
- Built a RAG-powered data catalogue using LangChain, LLMs, and semantic search.
- Containerised all services with Docker and automated deployments via CI/CD and Kubernetes.

**Data Engineer**, Analytics Ltd Aug 2016 – Mar 2019

- Migrated legacy MySQL pipelines to PostgreSQL, reducing query latency by 45%.
- Developed Python libraries for data validation used across 5 teams.

## Education

**MSc Data Science**, University of Alberta Sep 2014 – Jun 2016

## Projects

**StreamFlow** — Real-time IoT event processor built with Python, Kafka, and AWS.

- Ingests 200K events/minute; powers dashboards for model deployment monitoring.